



Hotpatching Embedded Real-time Applications

Patch your program without restarting it

Overview

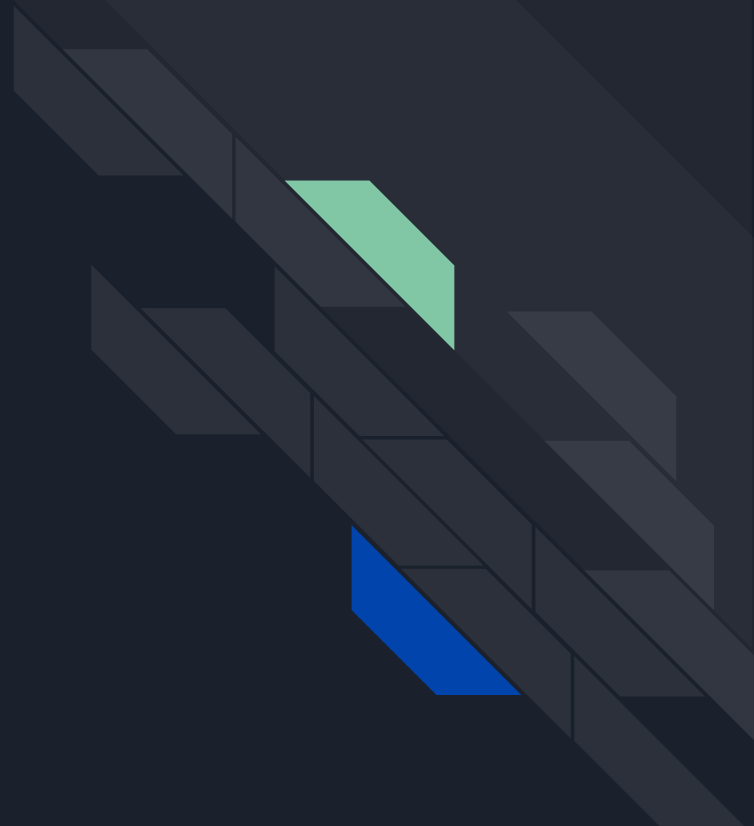
Motivation

Problem Statement & Challenges


Flash Patch & Breakpoint Unit

Patch Preparation & Switchover

Implementation



Motivation



- Updating vulnerable software - not possible always because of downtime and required reboot
 - Pacemaker, Plant Control IoT systems
- Patch updates introduce downtime
 - Scheduled patching windows
 - Exploit not-yet-updated software even when patches are available
- Hotpatching: update a program when it is being executed without stopping or restarting it



Challenges faced

- 01 No downtime. Continuous operation required. Cannot be stopped or restarted.
- 02 Changes in code execution must not violate real-time constraints. Small and predictable delay acceptable. Large and unpredictable are not.
- 03 Resource and Memory constraints.



Hotpatching Strategies

01

Hotpatching Relocatable Executables

Dynamically linked binaries reference code that is not part of the compiled library and is loaded during run-time. Dynamic linking requires only a symbolic link which can be adjusted during run-time to apply a hotpatch. In Embedded devices, binaries are statically linked. Hence not possible.

02

A/B Schemes

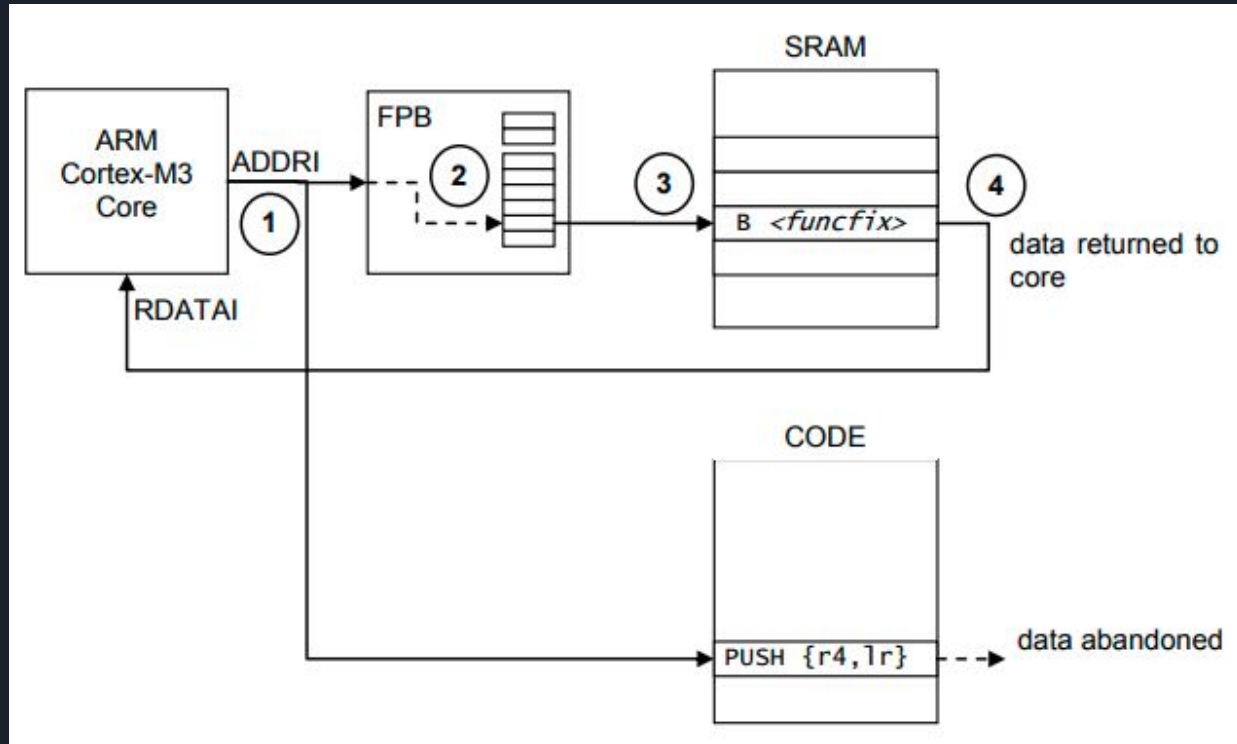
Two instances of firmware: One actively running while the other is updated. Switchover later. Problem: Doubling the memory to hold both instances. Dedicated firmware management unit required to update and switch.

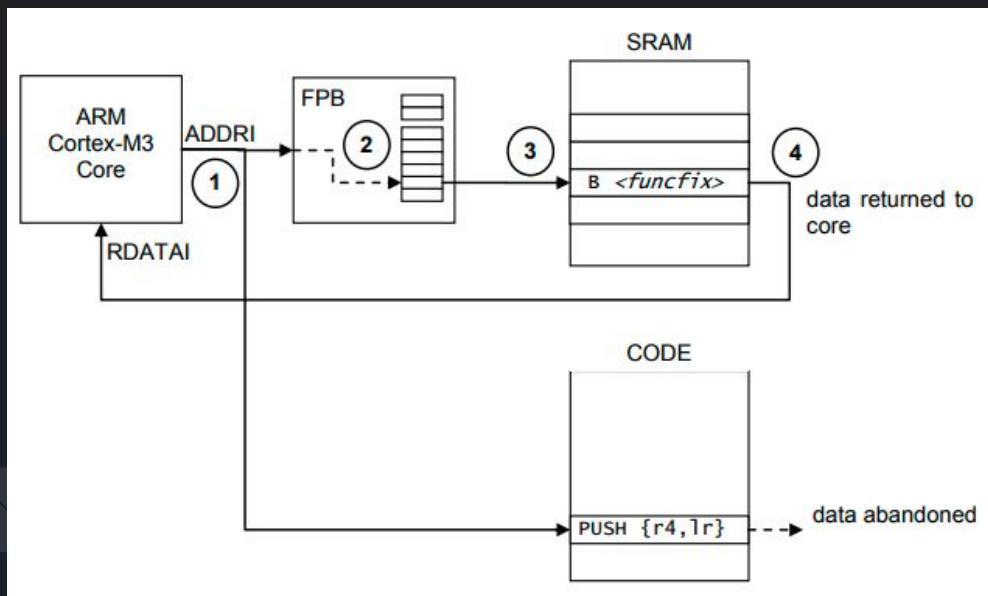
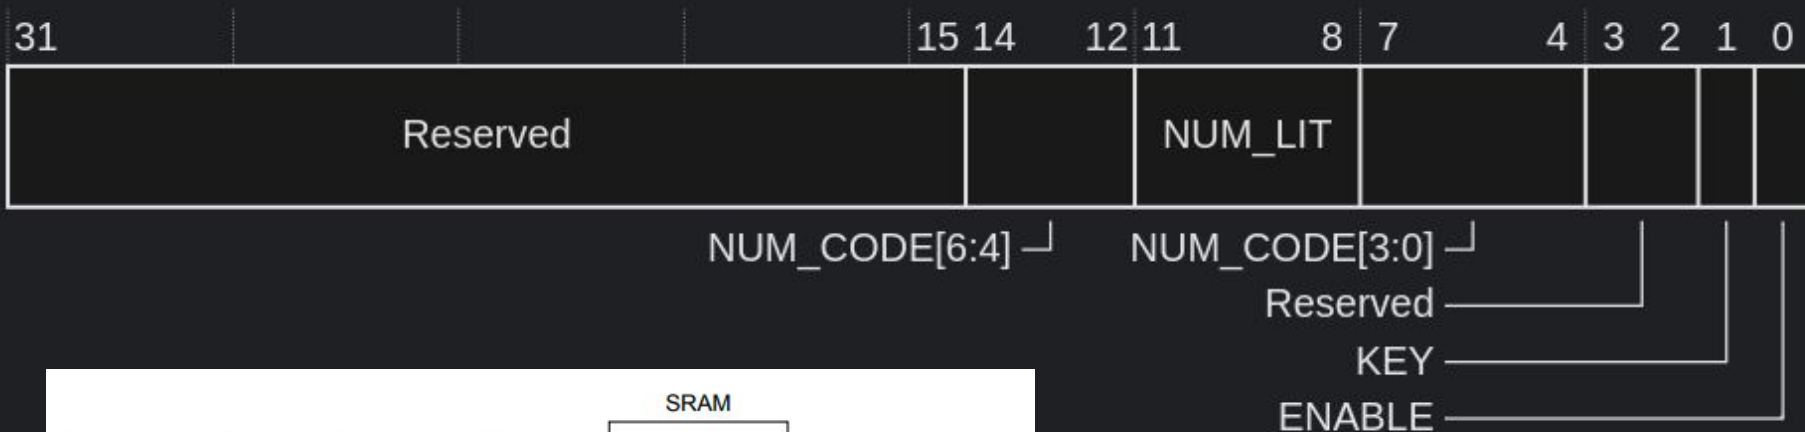


HERA: Key Idea

- Relies on standard debugging unit depending on the architecture
 - Onboard debugging unit with code-remapping unit required
 - temporarily stop the CPU, drop the fetched instruction, and fetch a new instruction
 - leverage this behavior to redirect the control flow to newly deployed hotpatches
 - Example: Flash Patch and Breakpoint in Cortex M3/M4 microprocessors
 - Compatible with existing hardware architectures
- Single trampoline instruction to activate the patch
 - Patch downloaded during IDLE times as per RTOS scheduling
 - Use code-remapping capability to execute a trampoline instead of a processor instruction
 - Trampoline jumps into the updated patch in a single instruction.
 - Execution time deterministic and can satisfy real-time constraints with hard deadlines
- Memory overheads minimal
 - Memory required only to store the patch
 - Couple of memory addresses


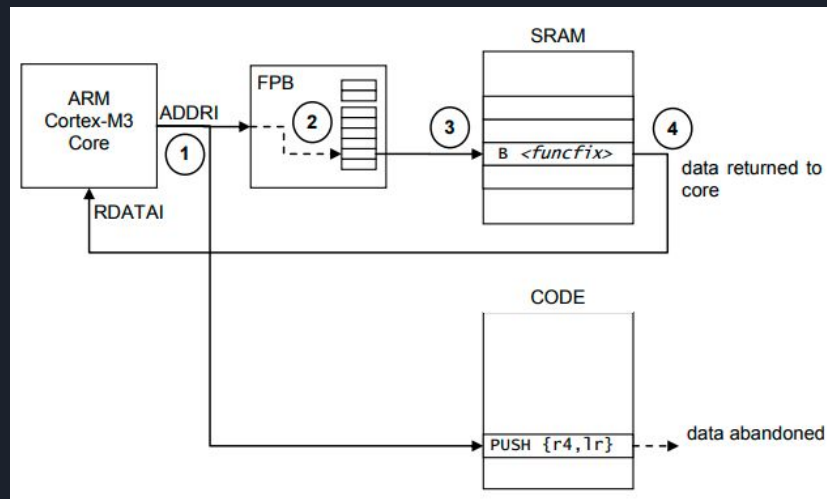
Flash Patch & Breakpoint Unit





FP_CTRL

Register Bit Information

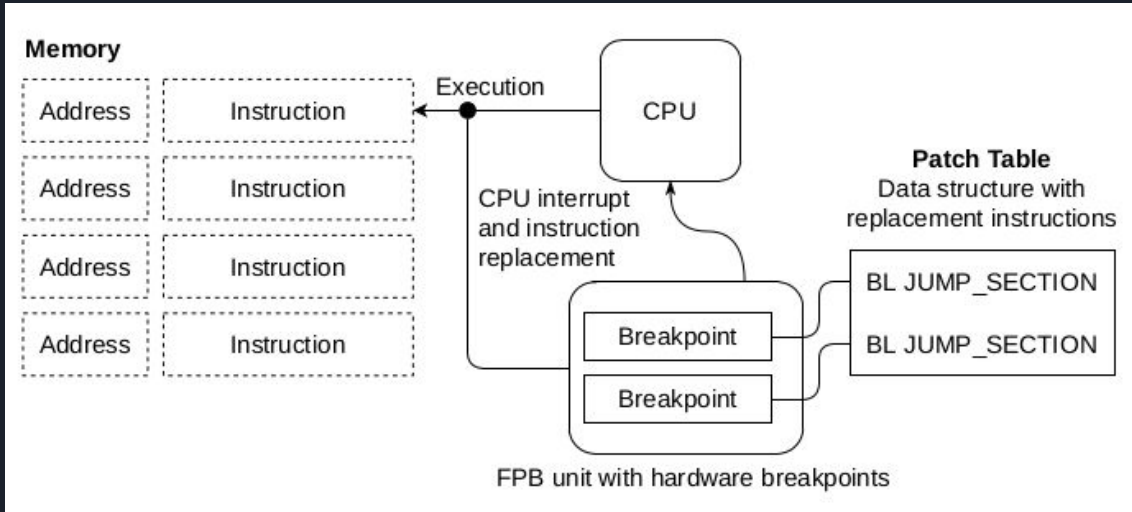


FP_REMAP
Register Bit
Information




FP_COMPx
Register Bit
Information

Flash Patch & Breakpoint Unit



Using FPB,

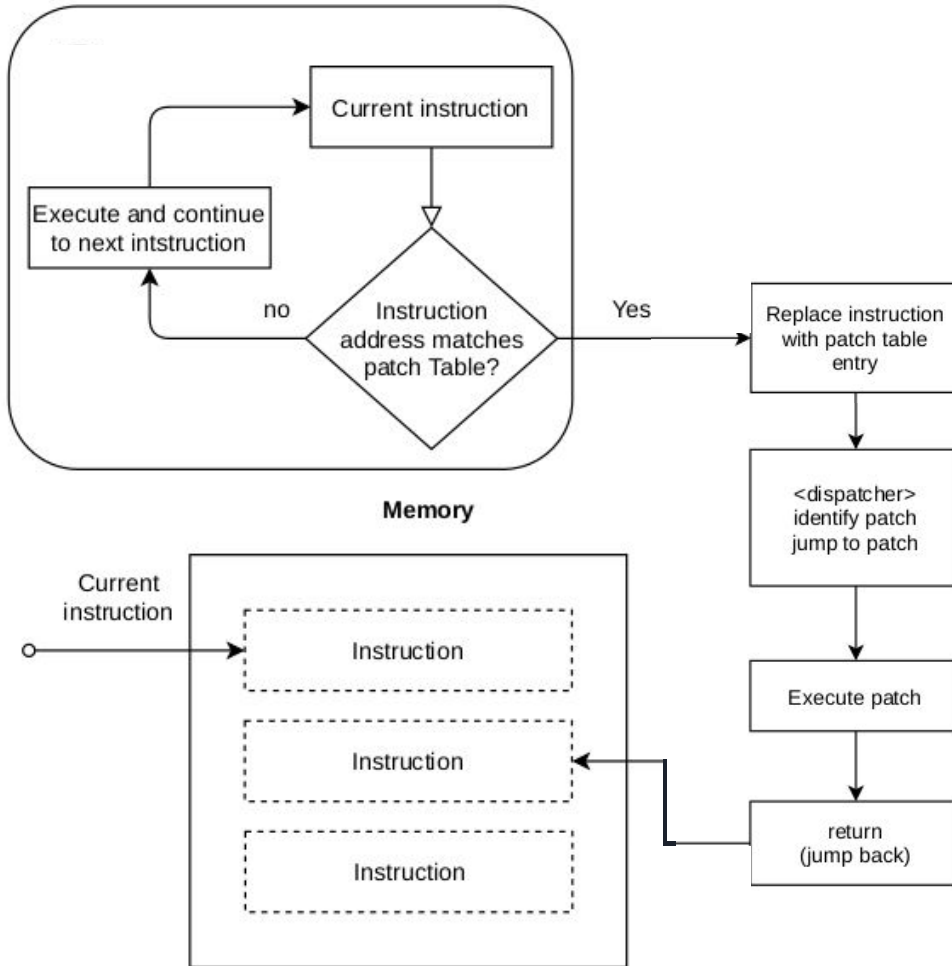
- Single 'BL' function call instruction can be replaced - leaving all other calls to the same function unaffected
- Replace an entire function: remap the first instruction. Note: Returning back is tricky



Patch Preparation

- Receive the patch from an external source,
- Verify the patch by checking its integrity and origin
- Parse the meta-information from the patch format
- Store the patch in selected memory area
- Add a new entry to the dispatcher,
- Configure the patching unit with insertion point and corresponding branching instruction (trampoline configuration)
- Trigger the atomic switchover.

Swichover



- Continuous monitoring of the currently executed instruction
- Instruction replacement on breakpoint hit
- Control-flow redirection to a trampoline
- Trampoline redirection to the appropriate patch
- Patch execution
- Return and continued execution.